

# Implementing Pub/Sub IPC in mCertikOS

## Understanding of Features

- **Topic Management**: Unique identifiers for message channels, with publishers sending messages to specific topics.
- **Subscriber Registration**: Processes can subscribe to topics, with optional message queuing (e.g., buffer size limits).
- **Message Broadcasting**: Publishers send messages to all active subscribers of a topic, ensuring non-blocking delivery.
- **Callback Mechanism**: Subscribers execute user-defined functions upon message receipt.
- **Resource Management**: Efficient allocation of memory for messages and queues, adhering to mCertikOS's container-based resource quotas

## Implementation Plan

### **Data Structure Design**

- **Topic Registry**: A hashmap mapping topic names to subscriber lists and message queues.
- **Subscriber Structure**: Each subscriber entry contains a process ID, callback function pointer, and a message queue.
- **Message Queue**: A circular buffer or linked list managed per subscriber, with configurable size (e.g., 1000 messages).

### **Layered Architecture**

#### **Layer 1: Message Management**

- **Functions**:
  - `message_create`: Allocates memory for a message (using `container_alloc`).
  - `message_enqueue`: Adds a message to a subscriber's queue, discarding old messages if the queue is full.

- `message_dequeue`: Retrieves the next message for processing.
- **Dependencies**:
  - **Container System**: Tracks memory usage per process (e.g., `container_can_consume` to prevent overallocation).
  - **Virtual Memory**: Uses page tables (`MPTOP` layer) for memory mapping.

## Layer 2: Topic Management

- **Functions**:
  - `topic_create`: Registers a new topic in the registry.
  - `topic_subscribe`: Adds a subscriber to a topic's list, initializing their message queue.
  - `topic_unsubscribe`: Removes a subscriber and frees associated resources.
- **Concurrency Control**: Use spin locks or semaphores from mCertikOS's synchronization primitives to handle concurrent subscriptions/unsubscriptions.

## Layer 3: Syscall Interface

- **New Syscalls**:
  - `sys_pub`: Publishes a message to a topic.
  - `sys_sub`: Subscribes to a topic, specifying a callback and queue size.
  - `sys_unsub`: Unsubscribes from a topic.
- **Integration**:
  - Modify the `syscall.h` header and `syscall.c` to include new syscalls.
  - Implement argument validation (e.g., topic existence, valid callback addresses).

## Layer 4: Callback Execution

- **Trap Handling**: Use mCertikOS's trap handling infrastructure (Lab 31) to schedule callback execution in user space.
  - When a message is enqueued, trigger a trap to the subscriber's process, invoking the callback.
  - Ensure callbacks execute in a non-blocking manner (e.g., via separate threads or interrupts).

# Key mCertikOS Functions to Leverage

## Memory Management

- **Container System:** Use `container_alloc/container_free` to manage memory for messages and queues, adhering to resource quotas.
- **Page Tables:** Utilize virtual memory layers (e.g., `MPTOp`, `MPTComm`) to map message data structures into user space.

## Process Management

- **Thread Creation:** Use existing process/thread management functions (e.g., `thread_create`) to handle concurrent callback execution.
- **Trap Handling:** Implement message delivery via trap handlers (e.g., `trap_init`, `trap_set_handler` from Lab 3).

## Synchronization Primitives

- **Locks:** Implement spin locks or semaphores (e.g., using x86 atomic operations) to protect shared data structures like topic registries.

## Roadmap for Implementation

### Phase 1: Core Data Structures and Memory Management

- **Define Topic Registry:** Implement a hash map for topics, storing subscriber lists and queues.
- **Implement Message Queues:** Use circular buffers or linked lists, managed via `container_alloc` for memory allocation.
- **Integrate Container System:** Ensure each process's message queue adheres to its resource quota<sup>1</sup>.

### Phase 2: Topic and Subscriber Management

- **Implement `topic_create`:** Add a new topic to the registry.
- **Implement `topic_subscribe`:**
  - Validate topic existence.
  - Allocate a message queue for the subscriber.
  - Store the subscriber's callback and queue size.
- **Implement `topic_unsubscribe`:**
  - Remove the subscriber from the topic's list.
  - Free the associated message queue and resources.

## Phase 3: Syscall Integration

- **Add Syscall Definitions:** Modify `syscall.h` to include `sys_pub`, `sys_sub`, and `sys_unsub`.
- **Implement Syscall Handlers:**
  - `sys_pub`: Retrieve the message, iterate over subscribers, and enqueue the message.
  - `sys_sub/sys_unsub`: Manage subscriber registrations.
- **Argument Validation:** Ensure topics exist before allowing operations.

## Phase 4: Message Delivery and Callbacks

- **Implement Message Broadcasting:**
  - For each subscriber, enqueue the message. If the queue is full, discard the oldest message.
- **Trigger Callback Execution:** Use mCertikOS's trap handling (Lab 31) to schedule the callback in user space.
  - When a message is enqueued, send an interrupt or trigger a trap to the subscriber.
  - Execute the callback in a non-blocking thread or via asynchronous traps.

## Technical Challenges and Mitigations

### Challenge 1: Atomic Operations for Shared Data

- **Solution:** Use x86 atomic instructions or spin locks to protect topic registries and queues during concurrent access.
- **Example:** Use `lock xadd` to increment subscriber counts atomically.

### Challenge 2: Message Delivery Deadlocks

- **Solution:** Ensure callbacks execute in a non-blocking manner. Use separate threads or prioritize message delivery via traps without blocking the CPU.

### Challenge 3: Memory Overallocation

- **Solution:** Strictly enforce resource quotas using mCertikOS's container system (`container_can_consume` before allocation1).